



A Brief Introduction to Reinforcement Learning

lisen.mu@hobot.cc



- **Recap and Concepts**
- Reinforcement Learning Basics
- Advanced Reinforcement Learning
- Challenges and Approaches



Recap and Concepts

- Machine Learning
- Supervised Learning
- Reinforcement Learning



Machine Learning?

- A computer program is said to learn:
 - from experience **E**
 - with respect to some class of tasks **T**
 - and performance measure **P**
- if its performance at tasks in T , as measured by P , improves with experience E

- Tom M. Mitchell

Supervised Learning Example:

Task A: Image Classification



Correctly associate labels with images: dog, plane, flower, cellphone, etc.

Supervised Learning Example:

Task A: Image Classification




T ask	Image Classification
E xperience	Labelled Image
P erformance Measure	Precision/Recall



Learning with Model

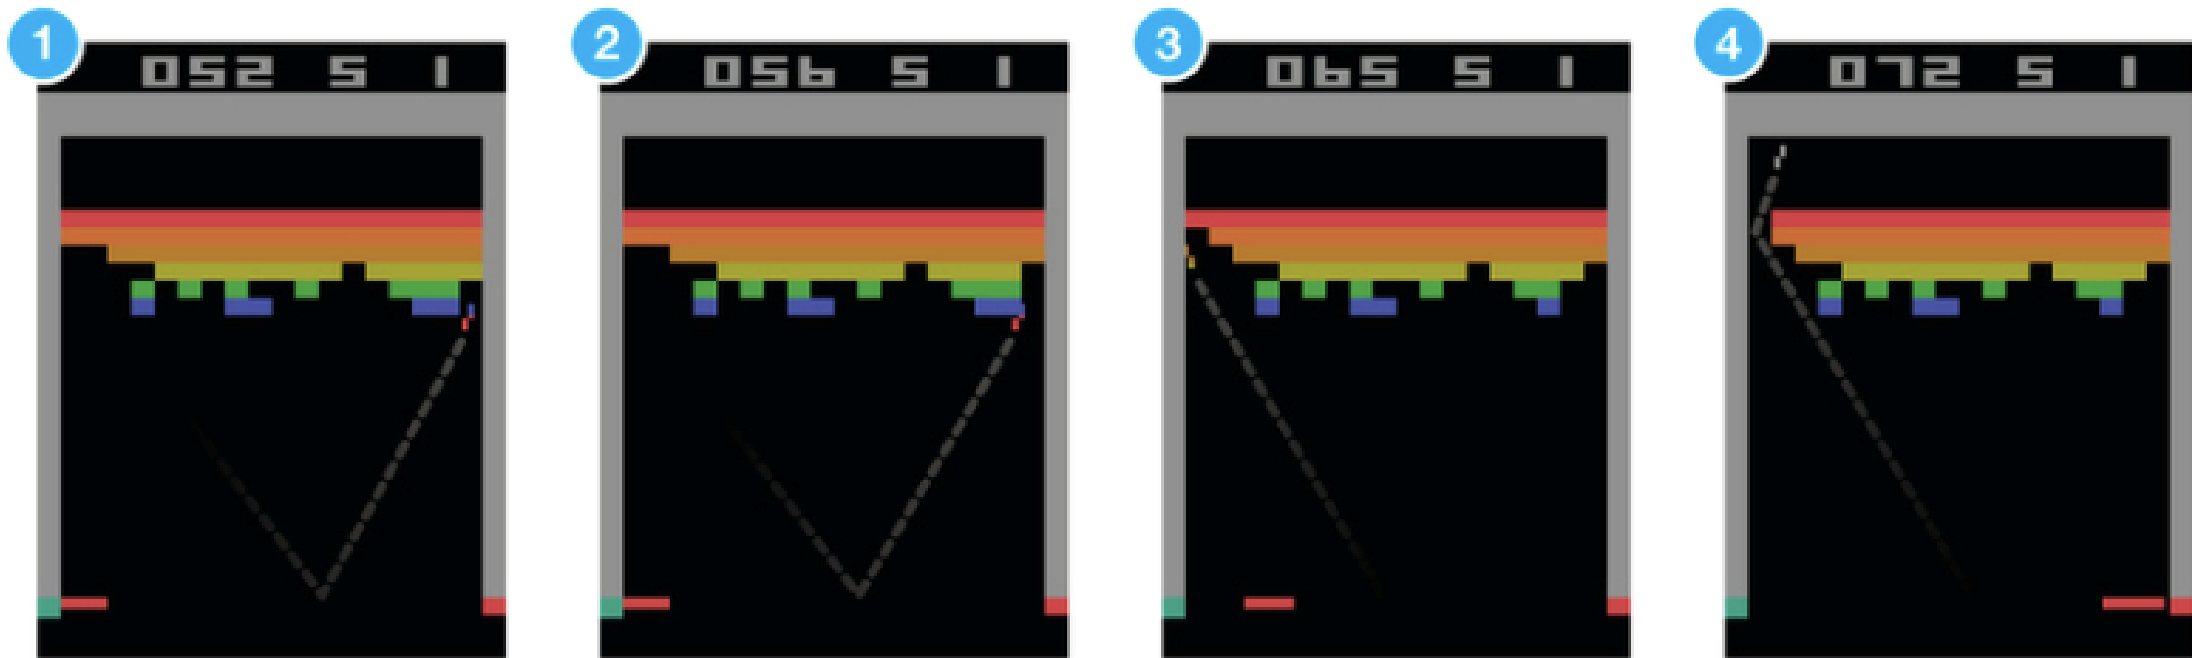
- True(unknown) function
 - *"how human classify objects in image?"*
- Hypothesis class(Model)
 - *function: $y = f(x)$ with parameters*
- Search within hypothesis space
 - according to dataset
 - with optimization methods

Learning with Model



Reinforcement Learning Example:

Task B: Play Game



Survive, and get highest score as possible

Reinforcement Learning Example:

Task B: Play Game



T ask	Play Game
E xperience	Play game over and over again
P erformance Measure	Overall score

Reinforcement Learning Example: Other Tasks





Reinforcement Learning

- How software **agents** ought to take **actions**
- in an **environment**, so as to maximize some notion of cumulative **reward**.



Reinforcement Learning

- Reinforcement learning is learning what to do
- how to map **situations** to **actions**
- so as to maximize a numerical **reward** signal.
- The learner is **not** told which actions to take
- but instead must **discover** which actions yield the most reward by trying them.

Reinforcement Learning



T ask	Optimal policy for specific task, under specific environment
E xperience	Interaction Experience with environment
P erformance Measure	Discounted Future Reward




T: Markov Decision Process

- $\langle \mathcal{S}, \mathcal{A}, \mathcal{R}, \mathcal{P} \rangle$
- \mathcal{S} : State Space
- \mathcal{A} : Action Space
- $\mathcal{R}: \mathcal{S} \times \mathcal{A} \rightarrow R$: Reward Function
- $\mathcal{P}: \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow [0, 1]$: Transition Function



E: Interaction Experience

- At each step t , the agent
 - receive state s_t
 - execute action a_t
- The environment
 - receive action a_t
 - emit scalar reward r_t
 - enter state s_{t+1}
- Transition: $\langle s_t, a_t, r_t, s_{t+1} \rangle$





P: Discounted Future Reward

- Total reward:

$$R = r_1 + r_2 + r_3 + \cdots + r_n$$

- Total future reward:

$$R_t = r_t + r_{t+1} + r_{t+2} + \cdots + r_n$$

- Discounted future reward:

$$R_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \cdots + \gamma^{n-t} r_n, \quad \gamma \in [0,1]$$



Policy, Value and Transition Model

- **Policy** is a behavior function choosing actions given states

$$a = \pi(s) \text{ or } p^{\pi}(a|s)$$

- **Value** function is expected discounted future award, starting from a given state and performing a given action

$$Q(s_t, a_t) = \mathbb{E}(R_t | s_t, a_t)$$

- Transition **model** estimates the future state(and reward)

$$p(s_{t+1}, r_t | s_t, a_t)$$



Approaches to RL

- **Value**-based RL
 - Estimate the optimal value function $Q^*(s, a)$
 - This is the maximum value achievable under any policy
- **Policy**-based RL
 - Search directly for the optimal policy π^*
 - This is the policy achieving maximum future reward
- **Model**-based RL
 - Build a transition model of the environment
 - Plan (e.g. by lookahead) using model $p(s_{t+1}, r_t | s_t, a_t)$



- Recap and Concepts
- **Reinforcement Learning Basics**
- Advanced Reinforcement Learning
- Challenges and Approaches



Reinforcement Learning Basics

- **Value Based RL**
- Policy Based RL
- Model Based RL



Q Function

- **Value** function is expected discounted future award, starting from a given state and performing a given action

$$Q(s_t, a_t) = \mathbb{E}(R_t | s_t, a_t)$$



Bellman Equation

- Optimal value function can be unrolled recursively

$$Q^*(s, a) = \mathbb{E}_{s'} \left(r + \gamma \max_{a'} Q^*(s', a') \mid s, a \right)$$

- Q-function can be iteratively updated by using the Bellman equation


$$Q_{i+1}(s, a) \leftarrow \mathbb{E}_{s'} \left(r + \gamma \max_{a'} Q_i(s', a') \mid s, a \right)$$



Example Problem: Grid World

- Q Function:


- $Q(s_0, a_l), Q(s_0, a_r), Q(s_1, a_l), Q(s_1, a_r), Q(s_3, a_l), Q(s_3, a_r)$
- $\gamma = 0.9$





Example Problem: Grid World

- $Q^*(s, a) = \mathbb{E}_{s'} \left(r + \gamma \max_{a'} Q^*(s', a') \mid s, a \right)$
- $Q(s_3, a_r), Q(s_1, a_l), Q(s_0, a_r), Q(s_1, a_r), Q(s_0, a_l), Q(s_3, a_l)$
- $\gamma = 0.9$



Value Based Method example: Q-Learning



```
initialize  $Q[num\_states, num\_actions]$  arbitrarily  
observe initial state  $s$   
repeat  
    select and carry out an action  $a$   
    observe reward  $r$  and new state  $s'$   
     $Q[s, a] = Q[s, a] + \alpha(r + \gamma \max_{a'} Q[s', a'] - Q[s, a])$   
     $s = s'$   
until terminated
```



Reinforcement Learning Basics

- Value Based RL
- **Policy Based RL**
- Model Based RL



Policy Function

- **Policy** is a behavior function choosing actions given states

$$a = \pi(s) \text{ or } p^{\pi}(a|s)$$


- Policy Based RL directly optimizes $\mathbb{E}[R]$ by searching in policy space



Example Problem: Grid World

- Policy Function:

- $P(a_l|s_0), P(a_r|s_0), P(a_l|s_1), P(a_r|s_1), P(a_l|s_3), P(a_r|s_3)$
- $\gamma = 0.9$






Example Problem: Grid World

- Policy Function:

- $P(a_l|s_0), P(a_r|s_0), P(a_l|s_1), P(a_r|s_1), P(a_l|s_3), P(a_r|s_3)$
- $\gamma = 0.9$



Policy Based Method example: Policy Gradient



- Define the loss function of policy $\pi(*: \theta)$ as

$$\mathcal{H}(\theta) = \mathbb{E}[R_t | \pi(*: \theta)] = \mathbb{E}[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots | \pi(*: \theta)]$$

- For stochastic policy $\pi(a|s: \mu)$

$$\frac{\partial \mathcal{H}(\theta)}{\partial \theta} = \mathbb{E} \left[Q^\pi(s, a) \frac{\partial \log \pi(a|s: \theta)}{\partial \theta} \right]$$

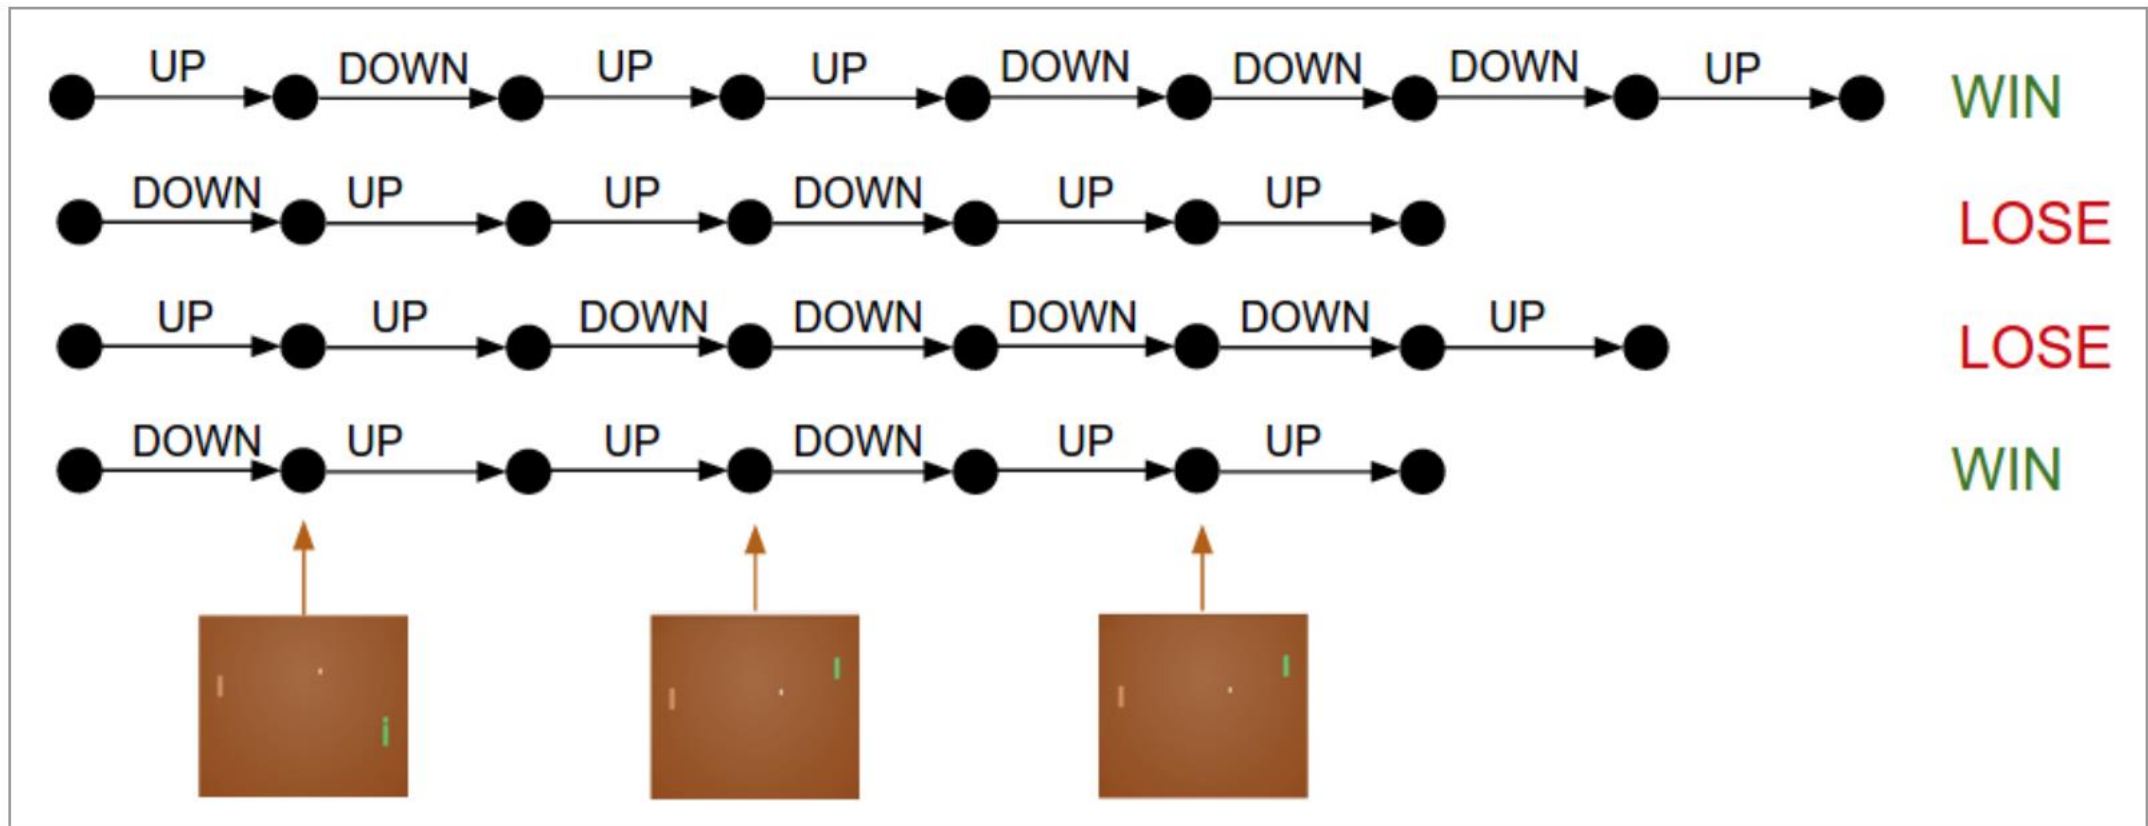
- For deterministic policy $a = \pi(s: \theta)$ when a is continuous and Q is differentiable

$$\frac{\partial \mathcal{H}(\theta)}{\partial \theta} = \mathbb{E} \left[\frac{\partial Q^\pi(s, a)}{\partial a} \frac{\partial a}{\partial \theta} \right]$$

Policy Based Method example: Policy Gradient



Policy Gradients: Run a policy for a while. See what actions led to high rewards. Increase their probability.





Reinforcement Learning Basics

- Value Based RL
- Policy Based RL
- **Model Based RL**



Model Based RL

- Transition **model** estimates the future state(and reward)

$$p(s_{t+1}, r_t | s_t, a_t), \text{ or } s_{t+1} = f(s_t, a_t)$$


- **Model**-based RL
 - Build a transition model of the environment
 - Plan (e.g. by lookahead) using model $p(s_{t+1}, r_t | s_t, a_t)$

Model Based RL Example: Model Predictive Control



1. run base policy $\pi_0(\mathbf{a}_t|\mathbf{s}_t)$ (e.g., random policy) to collect $\mathcal{D} = \{(\mathbf{s}, \mathbf{a}, \mathbf{s}')_i\}$
2. learn dynamics model $f(\mathbf{s}, \mathbf{a})$ to minimize $\sum_i \|f(\mathbf{s}_i, \mathbf{a}_i) - \mathbf{s}'_i\|^2$
3. plan through $f(\mathbf{s}, \mathbf{a})$ to choose actions
4. execute the first planned action, observe resulting state \mathbf{s}' (MPC)
5. append $(\mathbf{s}, \mathbf{a}, \mathbf{s}')$ to dataset \mathcal{D}

every N steps



Model Based RL Example: Model Predictive Control



1. run base policy $\pi_0(\mathbf{a}_t|\mathbf{s}_t)$ (e.g., random policy) to collect $\mathcal{D} = \{(\mathbf{s}, \mathbf{a}, \mathbf{s}')_i\}$
2. learn dynamics model $f(\mathbf{s}, \mathbf{a})$ to minimize $\sum_i \|f(\mathbf{s}_i, \mathbf{a}_i) - \mathbf{s}'_i\|^2$
3. plan through $f(\mathbf{s}, \mathbf{a})$ to choose actions **How?**
4. execute the first planned action, observe resulting state \mathbf{s}' (MPC)
5. append $(\mathbf{s}, \mathbf{a}, \mathbf{s}')$ to dataset \mathcal{D}

every N steps


Random, LQR, Backprop, etc. More on this later

Model Based RL Example: Model Predictive Control



1. run base policy $\pi_0(\mathbf{a}_t|\mathbf{s}_t)$ (e.g., random policy) to collect $\mathcal{D} = \{(\mathbf{s}, \mathbf{a}, \mathbf{s}')_i\}$
2. learn dynamics model $f(\mathbf{s}, \mathbf{a})$ to minimize $\sum_i \|f(\mathbf{s}_i, \mathbf{a}_i) - \mathbf{s}'_i\|^2$
3. plan through $f(\mathbf{s}, \mathbf{a})$ to choose actions
4. execute the first planned action, observe resulting state \mathbf{s}' (MPC)
5. append $(\mathbf{s}, \mathbf{a}, \mathbf{s}')$ to dataset \mathcal{D}

every N steps



Why?



REPLANNING HELPS WITH MODEL ERRORS

Model Based RL Example: Model Predictive Control



1. run base policy $\pi_0(\mathbf{a}_t|\mathbf{s}_t)$ (e.g., random policy) to collect $\mathcal{D} = \{(\mathbf{s}, \mathbf{a}, \mathbf{s}')_i\}$
2. learn dynamics model $f(\mathbf{s}, \mathbf{a})$ to minimize $\sum_i \|f(\mathbf{s}_i, \mathbf{a}_i) - \mathbf{s}'_i\|^2$
3. plan through $f(\mathbf{s}, \mathbf{a})$ to choose actions
4. execute the first planned action, observe resulting state \mathbf{s}' (MPC)
5. append $(\mathbf{s}, \mathbf{a}, \mathbf{s}')$ to dataset \mathcal{D} **Why?**

every N steps

Agent could go to states that π_0 would have never seen
Deal with distribution mismatch problem in \mathcal{D}



RL Basics: Recap

- **Value**-based RL
 - Estimate the optimal value function $Q^*(s, a)$
 - This is the maximum value achievable under any policy
- **Policy**-based RL
 - Search directly for the optimal policy π^*
 - This is the policy achieving maximum future reward
- **Model**-based RL
 - Build a transition model of the environment
 - Plan (e.g. by lookahead) using model $p(s_{t+1}, r_t | s_t, a_t)$



- Recap and Concepts
- Reinforcement Learning Basics
- **Advanced Reinforcement Learning**
- Challenges and Approaches

Advanced RL



- **DQN and Advanced Value Based RL**
- Advanced Policy Based RL
- Advanced Model Based RL






DQN and Advanced Value Based RL

- **DQN**
- Dueling DQN
- Double DQN
- Prioritized Experience Replay
- Optimality Tightening
- Rainbow

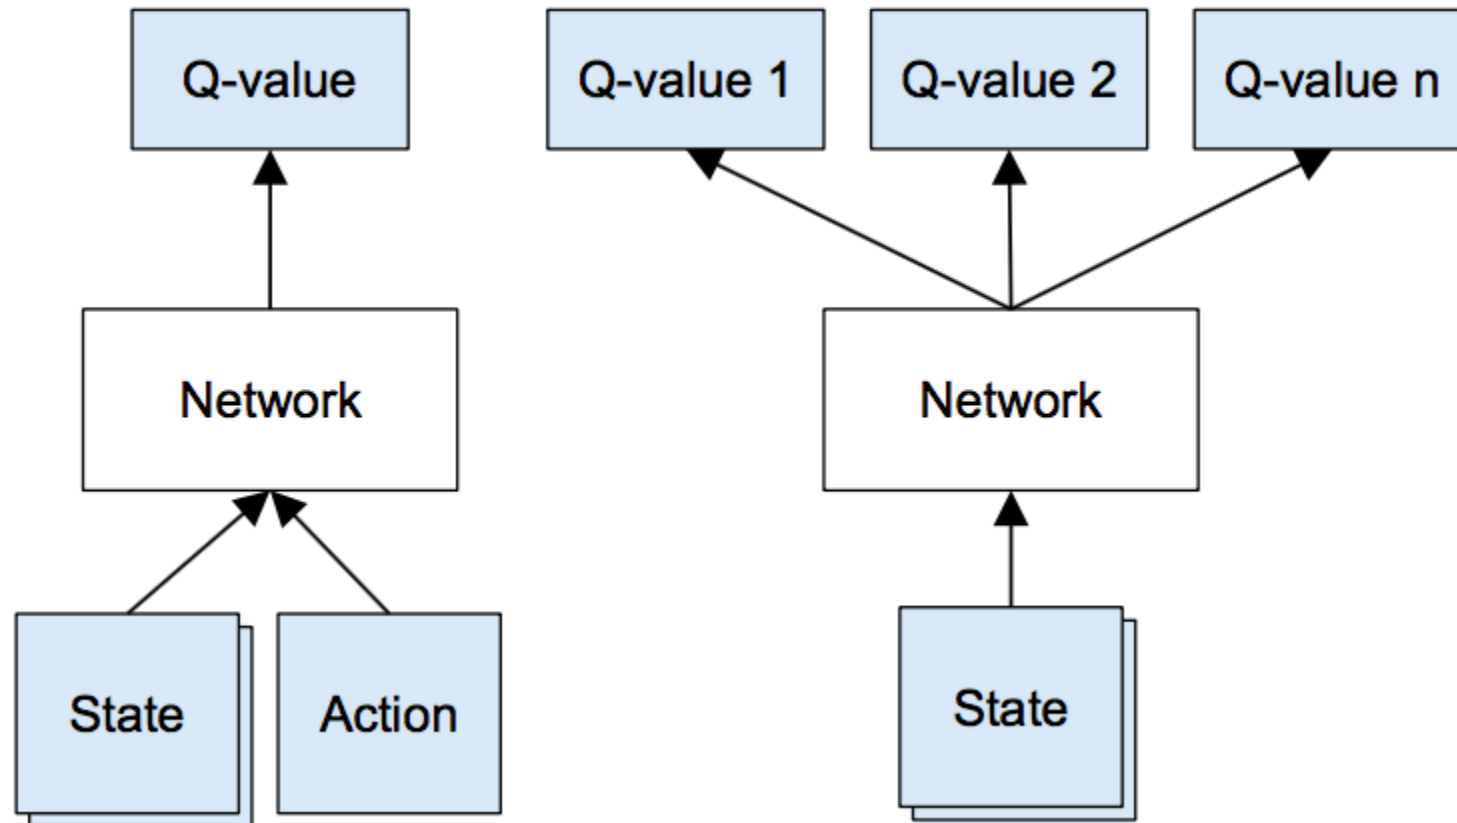
Recap: Neural Networks



- Universal Function Approximator
- Convolutional Network:
 - Function that accepts images as input



Deep Q-Network





Deep Q-Learning

- Represent value function by deep Q-network $Q(s, a: \theta)$
- Define a MSE loss function for Q-value approximation

$$\mathcal{L}(\theta) = \mathbb{E}_{s,a,r,s'} \left[\left(r + \gamma \max_{a'} Q(s', a': \theta) - Q(s, a: \theta) \right)^2 \right]$$

- Optimize by SGD

$$\frac{\partial \mathcal{L}(\theta)}{\partial \theta} = \mathbb{E}_{s,a,r,s'} \left[\left(r + \gamma \max_{a'} Q(s', a': \theta) - Q(s, a: \theta) \right) \frac{\partial Q(s, a: \theta)}{\partial \theta} \right]$$



Challenges in Deep Q-Learning

- Exploration-exploitation dilemma
 - Random exploration gradually becomes greedy and crude with exploitation of converging Q functions
- Sequential data are NOT i.i.d
 - Choosing an certain action affects the coming transitions
 - Highly correlated data are harmful for SGD method
- Non-stationary target
 - Target changes while θ is updated, causing oscillation during training
- Successful tricks:
 - ϵ -greedy exploration & experience replay & target Q-network

Deep Q-Learning with Experience Replay

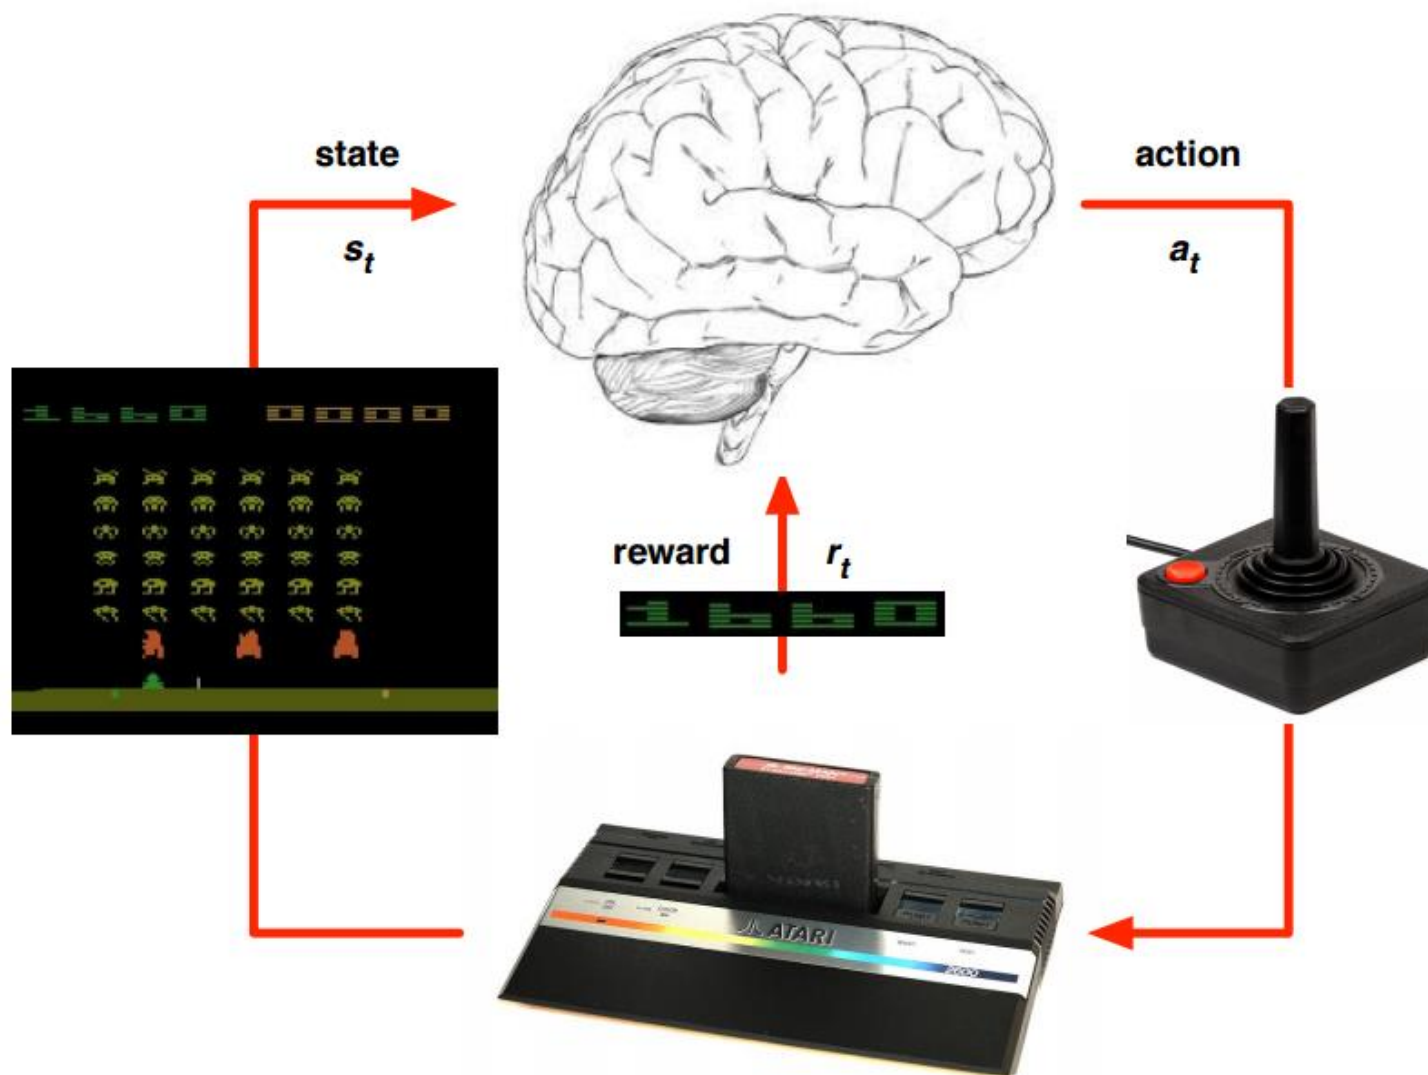


```
initialize replay memory  $D$ 
initialize action-value function  $Q$  with random weights
observe initial state  $s$ 
repeat
    select an action  $a$ 
        with probability  $\epsilon$  select a random action
        otherwise select  $a = \operatorname{argmax}_{a'} Q(s, a')$ 
    carry out action  $a$ 
    observe reward  $r$  and new state  $s'$ 
    store experience  $\langle s, a, r, s' \rangle$  in replay memory  $D$ 

    sample random transitions  $\langle ss, aa, rr, ss' \rangle$  from replay memory  $D$ 
    calculate target for each minibatch transition
        if  $ss'$  is terminal state then  $tt = rr$ 
        otherwise  $tt = rr + \gamma \max_{a'} Q(ss', aa')$ 
    train the  $Q$  network using  $(tt - Q(ss, aa))^2$  as loss

     $s = s'$ 
until terminated
```


RL in Atari






Deep Q-Learning in Atari

- End-to-end learning of values $Q(s, a)$ from pixels s
- Input state s is stack of raw pixels from last 4 frames
- Output is $Q(s, a)$ for 18 joystick/button positions
- Reward is change in score for that step



Results






DQN and Advanced Value Based RL

- DQN
- **Dueling DQN**
- Double DQN
- Prioritized Experience Replay
- Optimality Tightening
- Rainbow



Better Model

- Dueling Network Architectures for Deep Reinforcement Learning
- $Q(s, a; \theta, \alpha, \beta) = V(s; \theta, \beta) + A(s, a; \theta, \alpha)$





DQN and Advanced Value Based RL

- DQN
- Dueling DQN
- **Double DQN**
- **Prioritized Experience Replay**
- **Optimality Tightening**
- Rainbow



Convergence Problem in DQN

- Supervised Learning:

- $\mathcal{L}(\theta) = \mathbb{E}_x \left[(\mathbf{y} - f(x; \theta))^2 \right]$

- Reinforcement Learning(DQN):

- $\mathcal{L}(\theta) = \mathbb{E}_{s,a,r,s'} \left[\left(\mathbf{r} + \gamma \max_{a'} Q(s', a'; \theta) - Q(s, a; \theta) \right)^2 \right]$

- sparse ground truth

- Estimation bias in $\max_a Q$

- Only 1-step propagation along trajectory

Convergence Problem: Better use of data




- Prioritized Experience Replay
 - Some transitions are more informative than others
 - TD-error guided
 - for rarely visited valuable transitions

- $$P(i) = \frac{p_i^\alpha}{\sum_k p_k^\alpha}$$

- $$w(i) = \left(\frac{1}{N} * \frac{1}{P(i)}\right)^\beta$$


Recap: Bias and Variance of Value Estimation



Convergence Problem: Double DQN



- $Y^{DQN} = r + \gamma \max_{a'} Q(s', a' : \theta^-)$ **Upward bias!**
- $Y^{DoubleDQN} = r + \gamma Q(s', \max_{a'} Q(s', a' : \theta) : \theta^-)$
- Not too optimistic when estimating Q



Convergence Problem: Faster Q propagation



- Optimality Tightening
- $Q^*(s_t, a_t) = r + \gamma \max_a Q^*(s_{t+1}, a)$
 - $\geq \sum_{i=0}^k \gamma^i r_{t+i} + \gamma^{k+1} \max_a Q^*(s_{t+k+1}, a)$
- additional upper/lower bound along trajectory to Q

$$\min_{\theta} \sum_{(s_j, a_j, r_j, s_{j+1}) \in \mathcal{B}} \left[(Q_{\theta}(s_j, a_j) - y_j)^2 + \lambda (L_j^{\max} - Q_{\theta}(s_j, a_j))_+^2 + \lambda (Q_{\theta}(s_j, a_j) - U_j^{\min})_+^2 \right]$$




DQN and Advanced Value Based RL

- DQN
- Dueling DQN
- Double DQN
- Prioritized Experience Replay
- Optimality Tightening
- **Rainbow**

Rainbow



Combination of orthogonal improvements over DQN





Advanced RL

- DQN and Advanced Value Based RL
- **Advanced Policy Based RL**
- Advanced Model Based RL




Advanced Policy Based RL

- **Deep Policies**
- Reducing Variance in Policy Gradient Estimation
- Improved Sample Efficiency



Deep Policies

- Uses deep neural network as policy function approximators





Advanced Policy Based RL

- Deep Policies
- **Reducing Variance in Policy Gradient Estimation**
- Improved Sample Efficiency



Policy Gradient Recap

- **Policy** is a behavior function choosing actions given states

$$a = \pi(s) \quad \text{or} \quad p^\pi(a|s)$$

- Define the loss function of policy $\pi(*: \theta)$ as

$$\mathcal{H}(\theta) = \mathbb{E}[R_t | \pi(*: \theta)] = \mathbb{E}[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots | \pi(*: \theta)]$$

- For stochastic policy $\pi(a|s; \mu)$


$$\frac{\partial \mathcal{H}(\theta)}{\partial \theta} = \mathbb{E} \left[\boxed{Q^\pi(s, a)} \frac{\partial \log \pi(a|s; \theta)}{\partial \theta} \right]$$

How to estimate this?



Policy Gradient Estimation

- Estimate $\mathbb{E}[Q^\pi(s, a)]$ to get $\mathbb{E}\left[Q^\pi(s, a) \frac{\partial \log \pi(a|s; \theta)}{\partial \theta}\right]$
- Rollout after executing a !




$\mathbb{E}[R_i]$ is a good estimation, but very costly or not possible



Policy Gradient Estimation

- Estimate $\mathbb{E}[Q^\pi(s, a)]$ to get $\mathbb{E} \left[Q^\pi(s, a) \frac{\partial \log \pi(a|s; \theta)}{\partial \theta} \right]$
- Uses a single rollout




**R_0 is an unbiased estimation, but with high variance!
Further the horizon, higher the variance!**



Actor Critic

- Critic Function $V_{\phi}^{\pi}(s)$: Average R from state s under policy π
- $Q^{\pi}(s_t, a_t) = \sum_{i=0}^{k-1} \gamma^i r_{t+i} + \gamma^k V^{\pi}(s_{t+k})$
- $V_{\phi}^{\pi}(s)$: Bellman iteration



Shorter horizon(k), lower variance
But introduce bias with $V_{\phi}^{\pi}(s)$



Advantage v.s. $Q^\pi(s, a)$

- Policy Gradient: $\frac{\partial \mathcal{H}(\theta)}{\partial \theta} = \mathbb{E} \left[Q^\pi(s, a) \frac{\partial \log \pi(a|s;\theta)}{\partial \theta} \right]$
- $\mathbb{E} \left[(Q^\pi(s, a) - b(s)) \frac{\partial \log \pi(a|s;\theta)}{\partial \theta} \right]$ is also unbiased estimation
 - But possibly with lower variance
- $b(s)$: baseline function depending on state only
- Intuitively, the purpose of policy gradient is to evaluate relative goodness/badness of actions




Advantage based Actor Critic

- $V_{\phi}^{\pi}(s)$ is a good $b(s)$
- Advantage: $A^{\pi}(s, a) = Q^{\pi}(s, a) - V^{\pi}(s)$
- $\frac{\partial \mathcal{H}(\theta)}{\partial \theta} = \mathbb{E} \left[(Q^{\pi}(s, a) - b(s)) \frac{\partial \log \pi(a|s;\theta)}{\partial \theta} \right]$
 - $= \mathbb{E} \left[(\sum_{i=0}^{k-1} \gamma^i r_{t+i} + \gamma^k V^{\pi}(s_{t+k}) - V^{\pi}(s_t)) \frac{\partial \log \pi(a|s;\theta)}{\partial \theta} \right]$



Recap: Bias and Variance in Actor Critic


- $Q^\pi(s_t, a_t) = \sum_{i=0}^{k-1} \gamma^i r_{t+i} + \gamma^k V^\pi(s_{t+k})$
- Longer k , higher variance (from $\sum_{i=0}^{k-1} \gamma^i r_{t+i}$)
- Shorter k , higher bias (from $V^\pi(s_{t+k})$)





Generalized Advantage Estimation

- $Q^\pi(s_t, a_t) = \sum_{i=0}^{k-1} \gamma^i r_{t+i} + \gamma^k V^\pi(s_{t+k})$
- GAE: Weighted average of Q^π of different k
- A good balance between variance and bias






Advanced Policy Based RL

- Deep Policies
- Reducing Variance in Policy Gradient Estimation
- **Improved Sample Efficiency**



Policy Gradient: On-Policy Methods

- Policy Gradient: $\frac{\partial \mathcal{H}(\theta)}{\partial \theta} = \mathbb{E} \left[Q^\pi(s, a) \frac{\partial \log \pi(a|s;\theta)}{\partial \theta} \right]$
- On-Policy: Trajectory generated by π can only be used to estimate $Q^\pi(s, a)$ and update $\pi(\cdot | s; \theta)$




$\mathbb{E}[R_i]$ is a good estimation, but very costly!
Because every transition is used only once



Policy Gradient: On-Policy Methods

- Policy Gradient: $\frac{\partial \mathcal{H}(\theta)}{\partial \theta} = \mathbb{E} \left[Q^\pi(s, a) \frac{\partial \log \pi(a|s;\theta)}{\partial \theta} \right]$
- On-Policy: Trajectory generated by π can only be used to estimate $Q^\pi(s, a)$ and update $\pi(\cdot | s; \theta)$



Every transition is used only once



Policy Gradient: On-Policy Methods


- Can we do multiple policy update step based on a batch of transitions?
- Can we improve $\pi(\cdot | s; \theta_i)$ using transitions collected by $\pi(\cdot | s; \theta_0)$?



Trust Region Policy Optimization

- Monotonic policy improvement
 - Advantage between policies
 - Minorization-maximization
- Slightly off-policy
 - Trust Region: KL-divergence bounded
- Second order optimization



$$\begin{aligned} & \underset{\theta}{\text{maximize}} \mathbb{E}_{s \sim \rho_{\theta_{\text{old}}}, a \sim q} \left[\frac{\pi_{\theta}(a|s)}{q(a|s)} Q_{\theta_{\text{old}}}(s, a) \right] \\ & \text{subject to } \mathbb{E}_{s \sim \rho_{\theta_{\text{old}}}} [D_{\text{KL}}(\pi_{\theta_{\text{old}}}(\cdot|s) \parallel \pi_{\theta}(\cdot|s))] \leq \delta. \end{aligned}$$





Proximal Policy Optimization

- First order approximation to TRPO
 - KL divergence constraint in the form of loss term, instead of constrained optimization problem
- Easier to scale





Advanced RL

- DQN and Advanced Value Based RL
- Advanced Policy Based RL
- **Advanced Model Based RL**



Advanced Model Based RL

- **Planning Methods**
- Model Based with Model Free



Recap: Model Based RL

1. run base policy $\pi_0(\mathbf{a}_t|\mathbf{s}_t)$ (e.g., random policy) to collect $\mathcal{D} = \{(\mathbf{s}, \mathbf{a}, \mathbf{s}')_i\}$
2. learn dynamics model $f(\mathbf{s}, \mathbf{a})$ to minimize $\sum_i \|f(\mathbf{s}_i, \mathbf{a}_i) - \mathbf{s}'_i\|^2$
3. plan through $f(\mathbf{s}, \mathbf{a})$ to choose actions **How?**
4. execute the first planned action, observe resulting state \mathbf{s}' (MPC)
5. append $(\mathbf{s}, \mathbf{a}, \mathbf{s}')$ to dataset \mathcal{D}


every N steps

Random, LQR, Backprop, etc. More on this later



Tree Search Based Methods

- Monte Carlo tree search
 - Iteratively expand nodes into subtrees
- UCT Tree Search
 - Balance between promising nodes and underexplored nodes






LQR: Linear Quadratic Regulator

$$\min_{\mathbf{u}_1, \dots, \mathbf{u}_T} c(\mathbf{x}_1, \mathbf{u}_1) + c(f(\mathbf{x}_1, \mathbf{u}_1), \mathbf{u}_2) + \dots + c(f(f(\dots) \dots), \mathbf{u}_T)$$

$$c(\mathbf{x}_t, \mathbf{u}_t) = \frac{1}{2} \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix}^T \mathbf{C}_t \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix} + \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix}^T \mathbf{c}_t$$

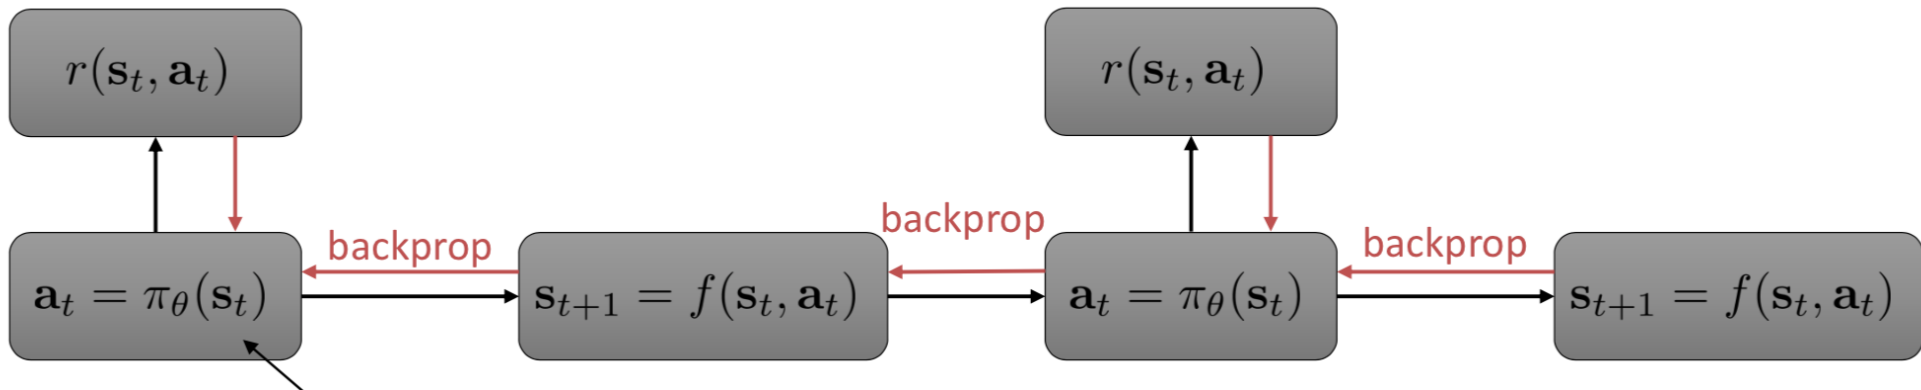
$$f(\mathbf{x}_t, \mathbf{u}_t) = \mathbf{F}_t \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix} + \mathbf{f}_t$$





Back-Propagation into Policy

- Backpropagate through model into the policy
- Choose action according to policy





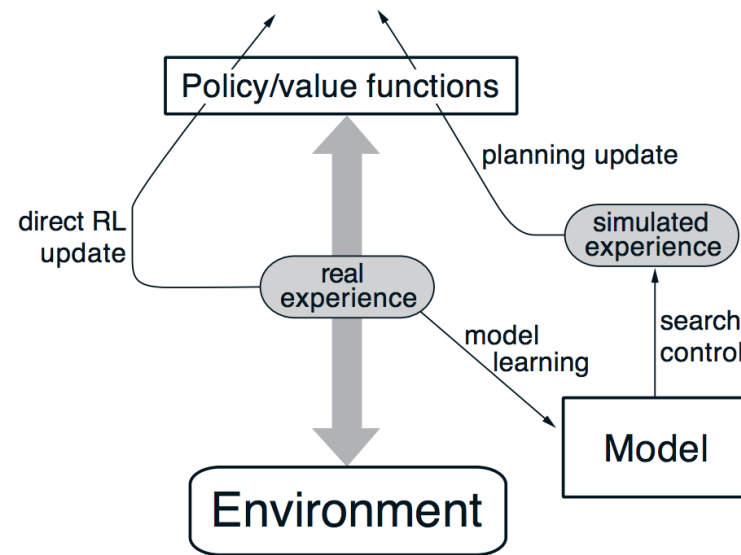
Advanced Model Based RL

- Planning Methods
- **Model Based with Model Free**



Dyna-Q: Model Based with Value Based

- Learn Q function and transition model
- Experience generate by transition model is also used to train Q function





Model Based with Policy Based


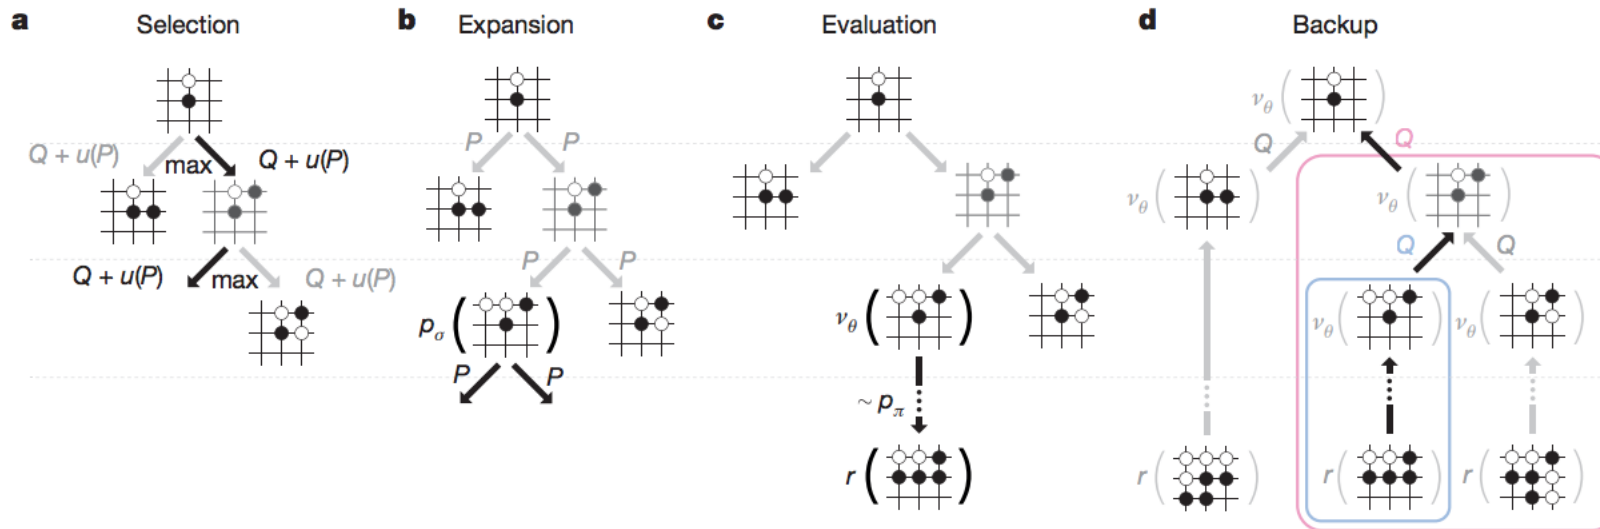
- Model Based DRL with Model-Free Fine-Tuning
 - Train model first
 - Use model(MPC) as expert to initialize Policy function
 - Policy Based RL afterwards

<https://arxiv.org/abs/1708.02596>

AlphaGo



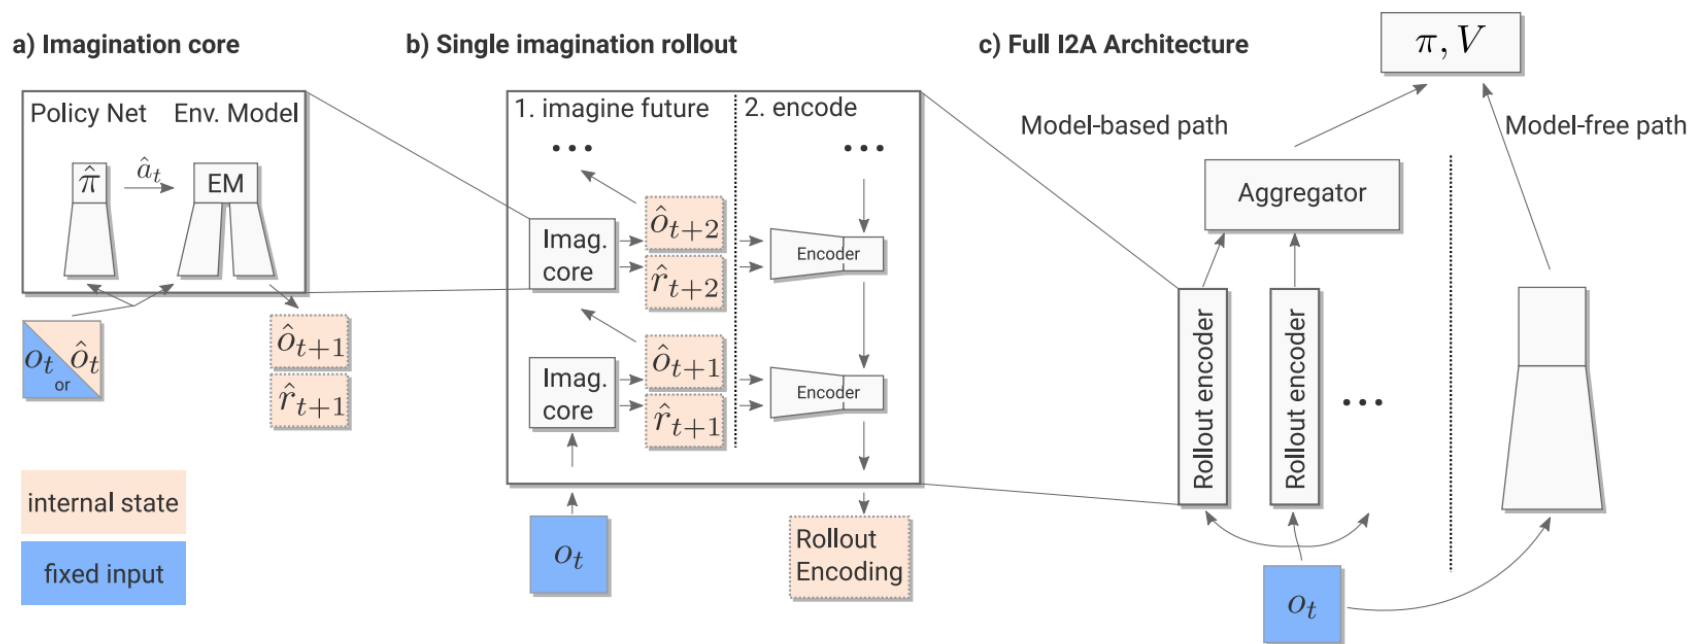
- Tree search into the future
- Value function to evaluate leaf node
- Policy function to choose moves





Imagination Augmented Agents

- Explicitly incorporates model prediction operation into policy network architecture



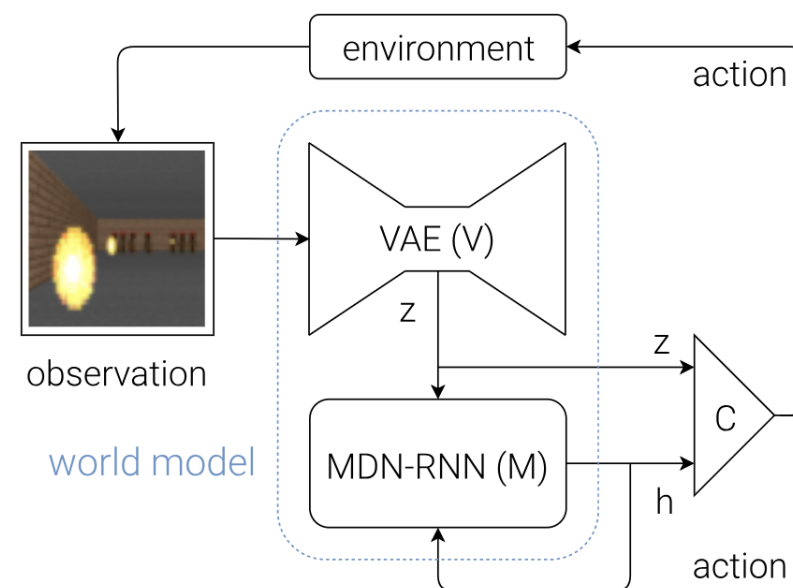


World Models

- Learn good latent state representation(with VAE)
- Train transition model on this latent state space
- Define & train simple controller(even with trajectory generated by transition model)



<https://arxiv.org/pdf/1803.10122>





- Recap and Concepts
- Reinforcement Learning Basics
- Advanced Reinforcement Learning
- **Challenges and Approaches**



Challenges and Approaches

- **Learning from Experts**
- Better Exploration
- Distributed Learning
- Hierarchical Methods
- Meta Learning



Learning From Experts

- Behavior cloning
 - Supervised Learning with data collected from experts
 - Cannot deal with states never seen before
- DAGGER
 - Query human experts for states unseen in previous dataset



Learning From Experts



- Inverse Reinforcement Learning
 - Recover reward function under which the experts trajectories are optimal
 - Iteratively optimize reward function and policy function



Learning From Experts

- Generative based Imitation Learning
 - Generative Adversarial Imitation Learning
 - DeepMimic

$$D_w : \mathcal{S} \times \mathcal{A} \rightarrow (0, 1)$$





Challenges and Approaches

- Learning from Experts
- **Better Exploration**
- Distributed Learning
- Hierarchical Methods
- Meta Learning



Better Exploration

- Exploration/Exploitation Dilemma
 - Exploit knowledge learnt to achieve higher reward
 - Explore new strategy to avoid local maximum




Better Exploration

- Noise in the Action Space
 - Q learning: Epsilon-Greedy
 - Policy: Entropy
 - $\lambda H(\pi; \theta)$
 - Deterministic Continuous Action: additional noise
 - $A = U + \mathcal{N}$
 - \mathcal{N} could be Gaussian or Ornstein-Uhlenbeck Noise



Better Exploration

- Noise in Policy Space
 - Bootstrapped DQN
 - Switching between heads for deep exploration
 - Stick to one head during entire episode
 - Noisy networks
 - Introduce noise in parameter space



Bootstrapped DQN

<https://arxiv.org/pdf/1602.04621.pdf>

<https://arxiv.org/pdf/1706.10295.pdf>

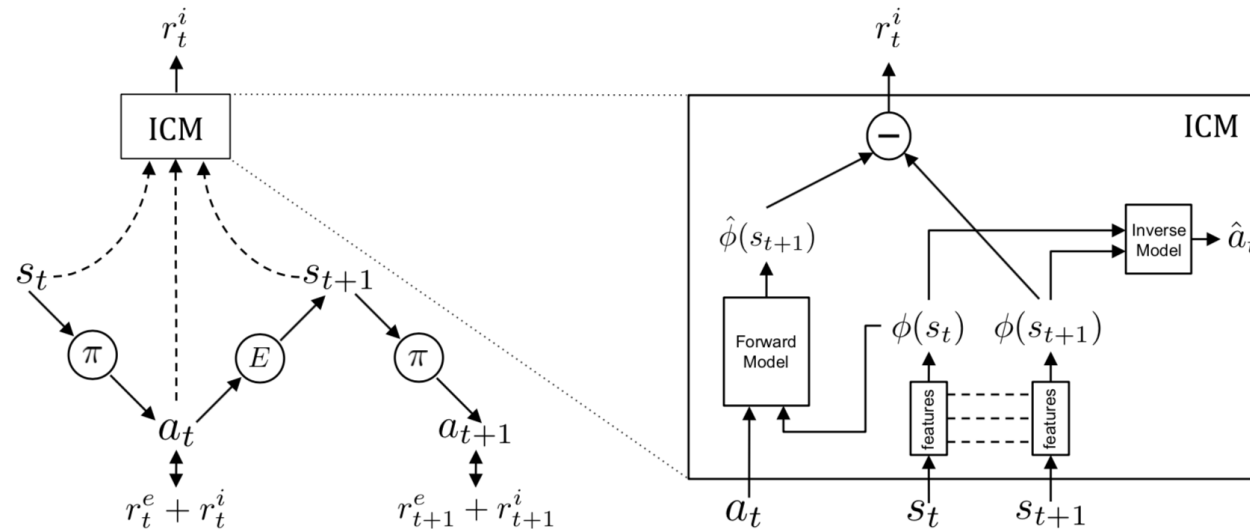


Better Exploration

- Intrinsic Reward

$$\mathcal{R}_{\text{Bonus}}(s, a) = \mathcal{R}(s, a) + \beta \mathcal{N}(s, a)$$


- Curiosity driven exploration: \mathcal{N} defined as loss of transition model






Better Exploration

- Learning to explore with meta-policy gradient
 - the policy gathering trajectories can be independent from the policy being optimized for the task
 - Meta reward: $\hat{\mathcal{R}}(D_0) = \hat{R}_{\pi'} - \hat{R}_{\pi}$



(d) Meta-Teacher (late)



(e) Meta-Student (late)



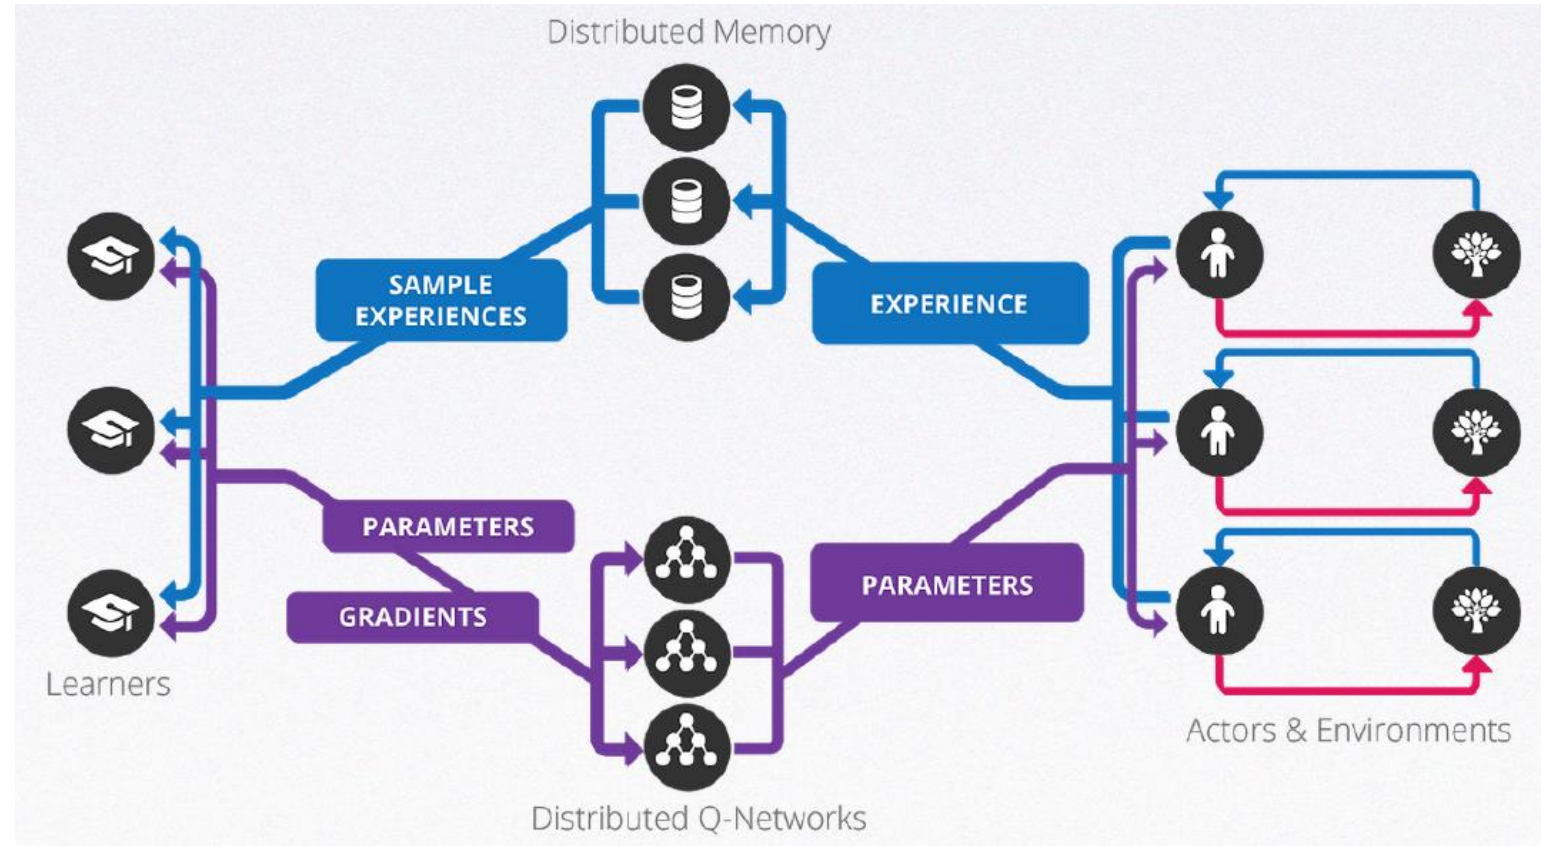
Challenges and Approaches

- Learning from Experts
- Better Exploration
- **Distributed Learning**
- Hierarchical Methods
- Meta Learning

Distributed Learning



Gorilla




Every component can be distributed



Distributed Learning

- Asynchronous methods for Deep RL





Distributed Learning

- Asynchronous methods for Deep RL
 - Superlinear speedup!

	Number of threads				
Method	1	2	4	8	16
1-step Q	1.0	3.0	6.3	13.3	24.1
1-step SARSA	1.0	2.8	5.9	13.1	22.1
n-step Q	1.0	2.7	5.9	10.7	17.2
A3C	1.0	2.1	3.7	6.9	12.5



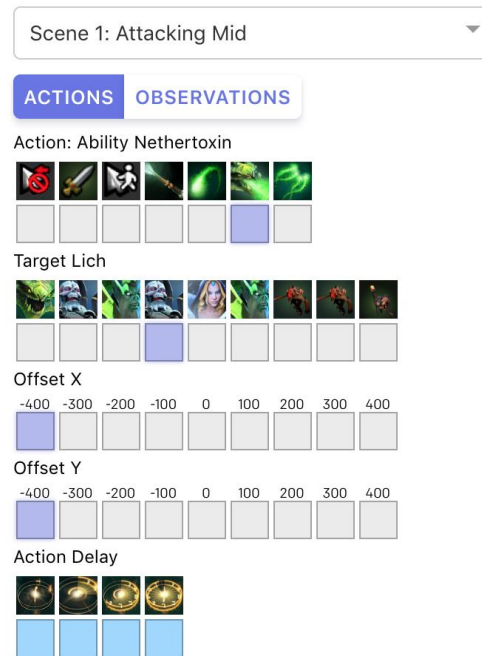
Distributed Learning

- AlphaGo Zero
 - 64 GPUs for optimization
 - 19 CPUs for parameter server

Distributed Learning



- OpenAI Five
 - PPO
 - 12800 CPUs
 - 256 GPUs
 - OpenAI Rapid





Challenges and Approaches

- Learning from Experts
- Better Exploration
- Distributed Learning
- **Hierarchical Methods**
- Meta Learning



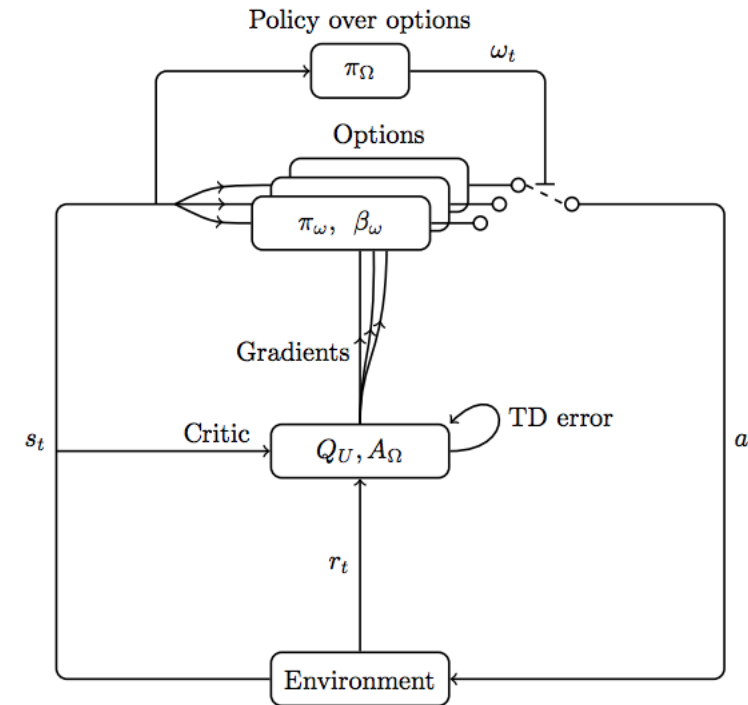
Hierarchical Methods

- Complex tasks require decisions to be made
 - At different time scale
 - At different abstraction level
- ...and can often be divided into simpler tasks handled by simpler policies

Hierarchical Methods



- Option Critic
 - Options defined as sub-policies
 - β_ω : termination condition
 - π_Ω : choose between options to complete task





Hierarchical Methods

- Universal Function Approximators
 - Universal Value Function: $Q^\pi(s, a, g)$
 - Universal Policy Function: $\pi(s||g)$
 - Hindsight Experience Replay
 - If the agent is given goal g but end up in g' , the agent still learns about how to achieve g'
- Scheduled Auxiliary Control
 - Make sure the agent receive very different sensory observation at multiple dimension



Challenges and Approaches

- Learning from Experts
- Better Exploration
- Distributed Learning
- Hierarchical Methods
- **Meta Learning**



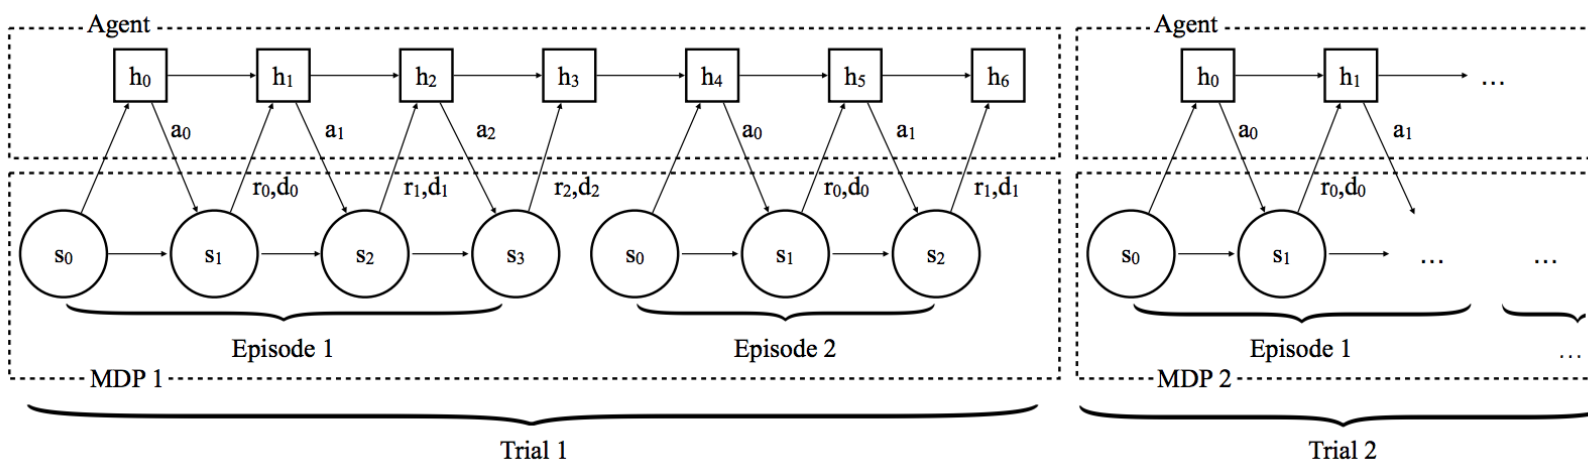
Meta Learning

- RL is slow(at least for now)
- What learnt from previous tasks would hopefully accelerate learning of future tasks



Meta Learning

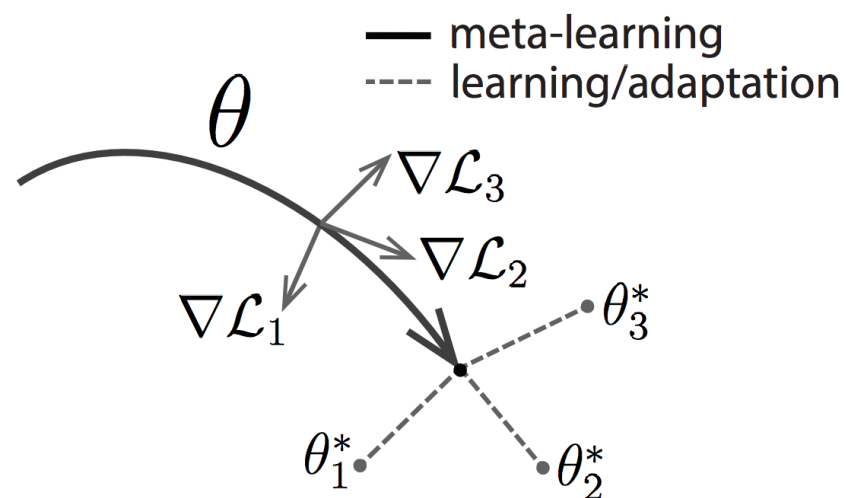
- Learning to Learn
 - RL^2 : Fast RL via Slow RL
 - The agent learns about how to solve new tasks as fast as possible





Meta Learning

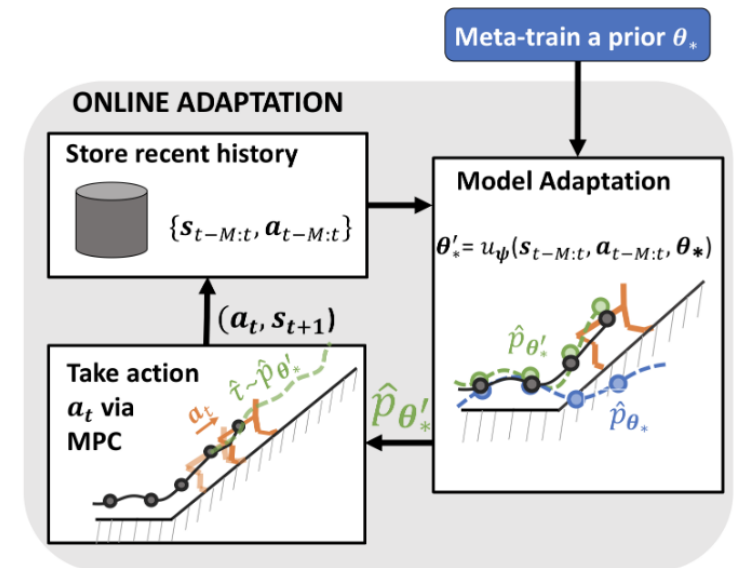
- Learning an initialization
 - Model-Agnostic Meta-Learning
 - Learns network parameter initialization for faster learning for a set of tasks



Meta Learning



- Learning an initialization
 - Learning to Adapt: Meta Learning for model based control
 - MAML for Model based RL
 - Learns about transition model parameter from recent interaction to the new environment





- Recap and Concepts
- Reinforcement Learning Basics
- Advanced Reinforcement Learning
- Challenges and Approaches

One More Thing



Practical Advices



Engineering is Very Important



Understand Your Task



Verify in Toy Tasks First



Look Closely into Data

Good Luck!



- Engineering is very important
- Understand your task
- Verify in toy tasks first
- Look closely into data

Q&A

Thank you.

And welcome:

dream@hobot.cc